

ML_with_SKLearn

November 6, 2022

1 Machine Learning in Python with SKLearn

By: Jonathan Blade

1.1 Reading the data in:

```
[14]: import pandas as pd

carData = pd.read_csv('/content/drive/MyDrive/Data/Auto.csv')
print(carData.head())
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70.0	
1	15.0	8	350.0	165	3693	11.5	70.0	
2	18.0	8	318.0	150	3436	11.0	70.0	
3	16.0	8	304.0	150	3433	12.0	70.0	
4	17.0	8	302.0	140	3449	NaN	70.0	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

1.2 Data Exploration:

```
[ ]: carData.mpg.describe()
```

```
[ ]: count    392.000000
      mean     23.445918
      std      7.805007
      min      9.000000
      25%     17.000000
      50%     22.750000
```

```
75%      29.000000
max      46.600000
Name: mpg, dtype: float64
```

```
[ ]: carData.year.describe()
```

```
[ ]: count      390.000000
      mean       76.010256
      std        3.668093
      min       70.000000
      25%       73.000000
      50%       76.000000
      75%       79.000000
      max       82.000000
      Name: year, dtype: float64
```

```
[ ]: carData.weight.describe()
```

```
[ ]: count      392.000000
      mean     2977.584184
      std      849.402560
      min     1613.000000
      25%     2225.250000
      50%     2803.500000
      75%     3614.750000
      max     5140.000000
      Name: weight, dtype: float64
```

The describe function tells us the ranges and averages of each column.

MPG: - Range: 9.0 - 46.0 - Average: 23.4

Year:

- Range: 70.0 - 82.0 - Average: 76.0

Weight:

- Range: 1613.0 - 5140.0 - Average: 2977.6

1.3 Analyzing Data Types:

```
[ ]: carData.dtypes
```

```
[ ]: mpg          float64
      cylinders    int64
      displacement float64
      horsepower   int64
      weight       int64
      acceleration float64
```

```
year          float64
origin        int64
name          object
dtype: object
```

```
[6]: carData.cylinders = carData.cylinders.astype('category').cat.codes
```

```
[5]: carData.origin = carData.origin.astype('category')
```

```
[7]: carData.dtypes
```

```
[7]: mpg          float64
cylinders        int8
displacement     float64
horsepower       int64
weight           int64
acceleration     float64
year             float64
origin           category
name            object
dtype: object
```

1.4 Dealing With N/A's:

```
[15]: carData.isnull().sum()
```

```
[15]: mpg          0
cylinders        0
displacement     0
horsepower       0
weight           0
acceleration     1
year             2
origin           0
name            0
dtype: int64
```

```
[16]: print('\nDimensions of data:', carData.shape)
carData = carData.dropna()
print('\nDimensions of data after removing N/A's:', carData.shape)
```

Dimensions of data: (392, 9)

Dimensions of data after removing N/A's: (389, 9)

1.5 Modifying Columns:

```
[44]: mpgAvg = carData.mpg.mean()
      mpgHigh = []

      counter = 0

      for x in carData.mpg:
          if x > mpgAvg:
              mpgHigh.append(1)
          else:
              mpgHigh.append(0)

      carData['mpg_high'] = mpgHigh

      carData = carData.drop('mpg', axis=1)

      print(carData.head())
```


	cylinders	displacement	horsepower	weight	acceleration	year	origin	\
0	8	307.0	130	3504	12.0	70.0	1	
1	8	350.0	165	3693	11.5	70.0	1	
2	8	318.0	150	3436	11.0	70.0	1	
3	8	304.0	150	3433	12.0	70.0	1	
6	8	454.0	220	4354	9.0	70.0	1	

	name	mpg_high
0	chevrolet chevelle malibu	0
1	buick skylark 320	0
2	plymouth satellite	0
3	amc rebel sst	0
6	chevrolet impala	0

1.6 Exploring the Data with Graphs:


```
[49]: import seaborn as sb

      graph = sb.catplot(x="mpg_high", kind='count', data=carData)
```




We can see from this graph that there are slightly more cars with below average MPG than there are cars with above average MPG.

```
[50]: graph2 = sb.relplot(x="horsepower", y="weight", data=carData)
```



We can determine from this graph that there appears to be a positive correlation between weight and horsepower.

```
[52]: graph3 = sb.boxplot(x='mpg_high', y='weight', data=carData)
```



This boxplot shows that cars with high MPG tend to weigh less than cars with low MPG.

1.7 Splitting the Data:

```
[54]: from sklearn.model_selection import train_test_split

X = carData.iloc[:, 0:7]
y = carData.iloc[:, 8]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=1234)

print('Dimensions of training data, X: ', X_train.shape, ' y: ', y_train.shape)
print('Dimensions of test data, X: ', X_test.shape, ' y: ', y_test.shape)
```

```
Dimensions of training data, X: (311, 7) y: (311,)
Dimensions of test data, X: (78, 7) y: (78,)
```

1.8 Logistic Regression:

```
[62]: from sklearn.linear_model import LogisticRegression

      clf1 = LogisticRegression(solver='lbfgs', max_iter=1000)
      clf1.fit(X_train, y_train)
      clf1.score(X_train, y_train)
```

```
[62]: 0.9067524115755627
```

```
[63]: pred1 = clf1.predict(X_test)
```

```
[64]: from sklearn.metrics import classification_report

      print(classification_report(y_test, pred1))
```

	precision	recall	f1-score	support
0	0.98	0.82	0.89	50
1	0.75	0.96	0.84	28
accuracy			0.87	78
macro avg	0.86	0.89	0.87	78
weighted avg	0.89	0.87	0.87	78

1.9 Decision Tree:

```
[65]: from sklearn.tree import DecisionTreeClassifier

      clf2 = DecisionTreeClassifier()
      clf2.fit(X_train, y_train)
```

```
[65]: DecisionTreeClassifier()
```

```
[66]: pred2 = clf2.predict(X_test)
```

```
[67]: print(classification_report(y_test, pred2))
```

	precision	recall	f1-score	support
0	0.92	0.90	0.91	50
1	0.83	0.86	0.84	28
accuracy			0.88	78
macro avg	0.87	0.88	0.88	78

weighted avg	0.89	0.88	0.89	78
--------------	------	------	------	----

```
[70]: from sklearn import tree
      tree.plot_tree(clf2)
```

```
[70]: [Text(0.6507352941176471, 0.9444444444444444, 'X[0] <= 5.5\ngini = 0.5\nsamples = 311\nvalue = [153, 158]'),
      Text(0.4338235294117647, 0.8333333333333334, 'X[2] <= 101.0\ngini = 0.239\nsamples = 173\nvalue = [24, 149]'),
      Text(0.27941176470588236, 0.7222222222222222, 'X[5] <= 75.5\ngini = 0.179\nsamples = 161\nvalue = [16, 145]'),
      Text(0.14705882352941177, 0.6111111111111112, 'X[1] <= 119.5\ngini = 0.362\nsamples = 59\nvalue = [14, 45]'),
      Text(0.058823529411764705, 0.5, 'X[4] <= 13.75\ngini = 0.159\nsamples = 46\nvalue = [4, 42]'),
      Text(0.029411764705882353, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
      Text(0.08823529411764706, 0.3888888888888889, 'X[3] <= 2683.0\ngini = 0.087\nsamples = 44\nvalue = [2, 42]'),
      Text(0.058823529411764705, 0.2777777777777778, 'X[3] <= 2377.0\ngini = 0.045\nsamples = 43\nvalue = [1, 42]'),
      Text(0.029411764705882353, 0.16666666666666666, 'gini = 0.0\nsamples = 38\nvalue = [0, 38]'),
      Text(0.08823529411764706, 0.16666666666666666, 'X[3] <= 2385.0\ngini = 0.32\nsamples = 5\nvalue = [1, 4]'),
      Text(0.058823529411764705, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
      Text(0.11764705882352941, 0.05555555555555555, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),
      Text(0.11764705882352941, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
      Text(0.23529411764705882, 0.5, 'X[3] <= 2567.0\ngini = 0.355\nsamples = 13\nvalue = [10, 3]'),
      Text(0.20588235294117646, 0.3888888888888889, 'X[5] <= 73.5\ngini = 0.469\nsamples = 8\nvalue = [5, 3]'),
      Text(0.17647058823529413, 0.2777777777777778, 'X[2] <= 88.0\ngini = 0.278\nsamples = 6\nvalue = [5, 1]'),
      Text(0.14705882352941177, 0.16666666666666666, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
      Text(0.20588235294117646, 0.16666666666666666, 'X[4] <= 17.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),
      Text(0.17647058823529413, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
      Text(0.23529411764705882, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
      Text(0.23529411764705882, 0.2777777777777778, 'gini = 0.0\nsamples = 2\nvalue =
```

```


[0, 2]'),
Text(0.2647058823529412, 0.3888888888888889, 'gini = 0.0\nsamples = 5\nvalue =
[5, 0]'),
Text(0.4117647058823529, 0.6111111111111112, 'X[3] <= 3250.0\ngini =
0.038\nsamples = 102\nvalue = [2, 100]'),
Text(0.35294117647058826, 0.5, 'X[3] <= 2880.0\ngini = 0.02\nsamples =
100\nvalue = [1, 99]'),
Text(0.3235294117647059, 0.3888888888888889, 'gini = 0.0\nsamples = 94\nvalue =
[0, 94]'),
Text(0.38235294117647056, 0.3888888888888889, 'X[3] <= 2920.0\ngini =
0.278\nsamples = 6\nvalue = [1, 5]'),
Text(0.35294117647058826, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue =
[1, 0]'),
Text(0.4117647058823529, 0.2777777777777778, 'gini = 0.0\nsamples = 5\nvalue =
[0, 5]'),
Text(0.47058823529411764, 0.5, 'X[2] <= 82.5\ngini = 0.5\nsamples = 2\nvalue =
[1, 1]'),
Text(0.4411764705882353, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.5, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.5882352941176471, 0.7222222222222222, 'X[4] <= 14.45\ngini =
0.444\nsamples = 12\nvalue = [8, 4]'),
Text(0.5588235294117647, 0.6111111111111112, 'X[5] <= 76.0\ngini =
0.444\nsamples = 6\nvalue = [2, 4]'),
Text(0.5294117647058824, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(0.5882352941176471, 0.5, 'X[3] <= 2760.0\ngini = 0.444\nsamples = 3\nvalue
= [2, 1]'),
Text(0.5588235294117647, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue =
[2, 0]'),
Text(0.6176470588235294, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.6176470588235294, 0.6111111111111112, 'gini = 0.0\nsamples = 6\nvalue =
[6, 0]'),
Text(0.8676470588235294, 0.8333333333333334, 'X[5] <= 79.5\ngini =
0.122\nsamples = 138\nvalue = [129, 9]'),
Text(0.7941176470588235, 0.7222222222222222, 'X[4] <= 21.6\ngini =
0.045\nsamples = 129\nvalue = [126, 3]'),
Text(0.7647058823529411, 0.6111111111111112, 'X[3] <= 2737.0\ngini =
0.031\nsamples = 128\nvalue = [126, 2]'),
Text(0.7058823529411765, 0.5, 'X[3] <= 2674.0\ngini = 0.444\nsamples = 3\nvalue
= [2, 1]'),
Text(0.6764705882352942, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue =
[2, 0]'),
Text(0.7352941176470589, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.8235294117647058, 0.5, 'X[2] <= 83.0\ngini = 0.016\nsamples = 125\nvalue
= [124, 1]'),

```

```

Text(0.7941176470588235, 0.3888888888888889, 'X[4] <= 18.55\ngini =
0.375\nsamples = 4\nvalue = [3, 1]'),
Text(0.7647058823529411, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.8235294117647058, 0.2777777777777778, 'gini = 0.0\nsamples = 3\nvalue =
[3, 0]'),
Text(0.8529411764705882, 0.3888888888888889, 'gini = 0.0\nsamples = 121\nvalue =
[121, 0]'),
Text(0.8235294117647058, 0.6111111111111112, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.9411764705882353, 0.7222222222222222, 'X[6] <= 1.5\ngini =
0.444\nsamples = 9\nvalue = [3, 6]'),
Text(0.9117647058823529, 0.6111111111111112, 'X[1] <= 247.0\ngini =
0.48\nsamples = 5\nvalue = [3, 2]'),
Text(0.8823529411764706, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(0.9411764705882353, 0.5, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(0.9705882352941176, 0.6111111111111112, 'gini = 0.0\nsamples = 4\nvalue =
[0, 4]')]]

```



1.10 Neural Network:

```

[71]: from sklearn import preprocessing

scaler = preprocessing.StandardScaler().fit(X_train)

X_scaledTrain = scaler.transform(X_train)

```

```
X_scaledTest = scaler.transform(X_test)
```

```
[72]: from sklearn.neural_network import MLPClassifier

      clf3 = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5,2), max_iter=500,
      ↪random_state=1234)
      clf3.fit(X_scaledTrain, y_train)
```

```
[72]: MLPClassifier(hidden_layer_sizes=(5, 2), max_iter=500, random_state=1234,
      solver='lbfgs')
```

```
[78]: pred3 = clf3.predict(X_scaledTest)
      print(classification_report(y_test, pred3))
```

	precision	recall	f1-score	support
0	0.94	0.88	0.91	50
1	0.81	0.89	0.85	28
accuracy			0.88	78
macro avg	0.87	0.89	0.88	78
weighted avg	0.89	0.88	0.89	78

```
[86]: clf4 = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(6,3,2), max_iter=500,
      ↪random_state=1234)
      clf4.fit(X_scaledTrain, y_train)
```

```
[86]: MLPClassifier(hidden_layer_sizes=(6, 3, 2), max_iter=500, random_state=1234,
      solver='lbfgs')
```

```
[87]: pred4 = clf4.predict(X_scaledTest)
      print(classification_report(y_test, pred4))
```

	precision	recall	f1-score	support
0	0.98	0.88	0.93	50
1	0.82	0.96	0.89	28
accuracy			0.91	78
macro avg	0.90	0.92	0.91	78
weighted avg	0.92	0.91	0.91	78

Both networks performed about the same, with the second network having slightly better results. I believe what led the second network to perform better was the additional hidden layer as well as a larger first hidden layer. However, one concern is overfitting the data which can occur with neural networks and small datasets.

1.11 Analysis:

The results of each algorithm are roughly comparable, with the second neural network competing just slightly better than the others.

Overall, precision for class 0 was over 0.9 for each algorithm with the highest being Logistic Regression and the second neural network, both with 0.98. However, the decision tree algorithm had the best values for both recall in class 0 and precision in class 1. The second neural network had the best recall for class 1, again tying with the logistic regression model. The second neural network manages to have the highest accuracy despite tying with the logistic regression model for both of its highest values.

Overall, the second neural network performed the best. This could be due to the neural network overfitting the data which becomes increasingly likely the smaller the dataset.

Overall, I think SKLearn feels easier to use and simpler than the equivalent functions in R. I still have a slight preference for R but possibly just because I am more familiar with it than SKLearn.